

Panoramic Image Stitching Using Feature-Based Registration and Blending

Digital Image Processing Project

Student Name: Mohanad Elbehairy , Ahmed Hany , Eyad Ahmed

Table of Contents

1. Abstract
 2. Introduction
 3. Theoretical Background
 4. System Architecture
 5. Algorithm Implementation
 6. Experimental Results
 7. Conclusion
 8. References
-

Abstract

This project implements a robust panoramic image stitching system that combines multiple overlapping images into a seamless wide-field panorama. The system employs feature-based image registration using SIFT and ORB descriptors, robust homography estimation with RANSAC, and implements novel sanity checking mechanisms to prevent common failure modes. Additional contributions include cylindrical projection for wide-angle correction, affine transformation fallback for stability, and memory-safe canvas management. The system achieves a 99% success rate with an average inlier ratio of 76% and processes typical 4-image panoramas in under 6 seconds.

1. Introduction

1.1 Problem Statement

Panoramic image stitching addresses the challenge of combining multiple overlapping photographs into a single coherent wide-angle image. Given a sequence of n images $\{I_1, I_2, \dots, I_n\}$ captured from different viewpoints, the objective is to:

- Identify corresponding features between adjacent images
- Estimate geometric transformations that align the images
- Warp images into a common reference frame
- Blend overlapping regions to create seamless transitions

1.2 Applications

- Landscape and architectural photography
- Virtual reality and 360° tours
- Medical imaging (X-ray panoramas)
- Satellite and aerial imagery mosaicking
- Document scanning and restoration

1.3 Challenges

- **Geometric alignment:** Images captured from different angles require perspective transformation
- **Illumination variation:** Lighting changes between captures create visible seams
- **Outlier matches:** False feature correspondences corrupt transformation estimation
- **Computational efficiency:** Processing high-resolution images in reasonable time
- **Memory constraints:** Large panoramas can exceed available RAM

1.4 Project Objectives

This implementation focuses on:

1. Robust feature detection and matching
2. Reliable geometric estimation with failure recovery
3. High-quality seamless blending
4. Interactive user interface for experimentation

2. Theoretical Background

2.1 Feature Detection

2.1.1 What is a Feature?

A feature is a distinctive point in an image that can be reliably detected across different views. Good features exhibit:

- **Repeatability:** Detectable in multiple images of the same scene
- **Distinctiveness:** Unique descriptor for accurate matching
- **Locality:** Robust to occlusions and clutter
- **Efficiency:** Fast to compute and compare

2.1.2 SIFT (Scale-Invariant Feature Transform)

SIFT detects and describes features that are invariant to scale, rotation, and partially invariant to illumination and viewpoint changes.

Detection Process:

1. **Scale-space construction:** Convolve image with Gaussian kernels at multiple scales
2. $L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$
3. **Keypoint localization:** Find extrema in Difference of Gaussians (DoG)
4. $D(x, y, \sigma) = L(x, y, k\sigma) - L(x, y, \sigma)$
5. **Orientation assignment:** Compute gradient orientation histogram
6. **Descriptor generation:** 128-dimensional vector from gradient magnitudes and orientations

Properties:

- Highly distinctive descriptors
- Robust to scale and rotation
- Computationally expensive (~300ms per image)

2.1.3 ORB (Oriented FAST and Rotated BRIEF)

ORB is a fast alternative using binary descriptors.

Detection Process:

1. **FAST corner detection:** Identify corners by comparing pixel intensities

2. **Orientation via intensity centroid:** Compute dominant direction
3. **BRIEF descriptors:** 256-bit binary string from intensity comparisons

Properties:

- $100\times$ faster than SIFT
- Binary descriptors enable Hamming distance matching
- Less distinctive than SIFT but sufficient for most panoramas

2.2 Feature Matching

2.2.1 Brute Force k-NN Matching

For each descriptor in image 1, find the $k=2$ nearest neighbors in image 2 based on descriptor distance:

- **SIFT:** Euclidean distance (L2 norm)
- **ORB:** Hamming distance (XOR + bit count)

2.2.2 Lowe's Ratio Test

Filters ambiguous matches by comparing distances to first and second nearest neighbors:

Accept match if: $d_1/d_2 < \text{ratio_threshold}$

where:

d_1 = distance to nearest neighbor

d_2 = distance to second nearest neighbor

ratio_threshold = 0.75 (SIFT) or 0.70 (ORB)

Intuition: If the best match is much closer than the second-best, it's likely correct. If they're similar, the match is ambiguous.

Effect: Typically removes 70-80% of initial matches, retaining only distinctive correspondences.

2.2.3 Mutual Matching (Bidirectional Consistency)

Additional filtering requiring consistency in both matching directions:

Keep match (i,j) only if:

- Feature i in image A matches to feature j in image B
- Feature j in image B matches back to feature i in image A

Implementation:

```
# Forward matching A → B
matches_AB = knn_match(descriptors_A, descriptors_B, k=2)

# Backward matching B → A
matches_BA = knn_match(descriptors_B, descriptors_A, k=2)

# Keep mutual matches
mutual = intersection(matches_AB, matches_BA)
```

Effect: Reduces false positives by 60-70%, significantly improving transformation estimation quality.

2.3 Homography Estimation

2.3.1 Homography Fundamentals

A homography H is a 3×3 matrix representing a planar projective transformation:

$$[x'] = [h_{11} \ h_{12} \ h_{13}] [x]$$

$$[y'] = [h_{21} \ h_{22} \ h_{23}] \times [y]$$

$$[w'] = [h_{31} \ h_{32} \ 1] [1]$$

Final coordinates: $x' = x'/w'$, $y' = y'/w'$

Degrees of freedom: 8 (9 elements, but scale is arbitrary)

When valid:

- Camera undergoes pure rotation (panorama case)
- Scene is approximately planar
- Combination of rotation and planar scene

2.3.2 Direct Linear Transformation (DLT)

Estimates H from ≥ 4 point correspondences. Each correspondence $(x_i, y_i) \leftrightarrow (x'_i, y'_i)$ provides 2 equations:

$$h_{11}x_i + h_{12}y_i + h_{13} - h_{31}x'_i x_i - h_{32}x'_i y_i - x'_i = 0$$

$$h_{21}x_i + h_{22}y_i + h_{23} - h_{31}y'_i x_i - h_{32}y'_i y_i - y'_i = 0$$

This forms a system $\mathbf{Ah} = \mathbf{0}$ solved via Singular Value Decomposition (SVD).

2.3.3 RANSAC (RANDOM SAMPLING AND consensus)

Feature matching inevitably produces outliers. RANSAC robustly estimates models in the presence of outliers.

Algorithm:

Input: Point correspondences, threshold t, iterations N

Output: Best homography H, inlier set

1. FOR i = 1 to N:

a. Randomly select 4 point correspondences

b. Compute homography H using DLT

c. For each correspondence j:

- Compute error: $e = ||x'_j - H \cdot x_j||$

- If $e < t$: mark as inlier

d. Count inliers

2. Return H with maximum inliers

3. Optional: Refine H using all inliers

Parameters:

- Reprojection threshold: $t = 4$ pixels (SIFT), 6 pixels (ORB)
- Iterations: Adaptive based on inlier ratio
- Success probability: 99%

USAC/MAGSAC: Modern variants with improved convergence and adaptive thresholding.
Used when available.

2.4 Homography Sanity Checking

2.4.1 The Problem

RANSAC can produce mathematically valid but geometrically implausible homographies, causing:

- Extreme perspective distortion
- Massive projected image sizes ($>10\times$)
- Characteristic "black wedge" artifacts
- Memory overflow

2.4.2 Sanity Criteria

Test 1: Perspective Constraint

Reject if: $|h_{31}| > 0.01$ OR $|h_{32}| > 0.01$

Rationale: Large values indicate extreme out-of-plane rotation, unlikely in panorama capture.

Test 2: Size Constraint

Project four corners of source image through H and compute bounding box:

corners = [(0,0), (w,0), (w,h), (0,h)]

projected = H \times corners

bbox = bounding_box(projected)

Reject if:

bbox.width > 3.0 \times w OR

bbox.height > 3.0 \times h

Rationale: 3 \times expansion is already generous. Larger projections indicate incorrect matches.

Implementation:

```
def homography_is_sane(H, img_shape, max_expand=3.0):  
  
    # Test 1: Perspective  
  
    if abs(H[2,0]) > 0.01 or abs(H[2,1]) > 0.01:  
  
        return False
```

```

# Test 2: Size

h, w = img_shape[:2]

corners = np.float32([[0,0], [w,0], [w,h], [0,h]])

projected = cv2.perspectiveTransform(corners, H)

w_proj = projected[:,0].max() - projected[:,0].min()

h_proj = projected[:,1].max() - projected[:,1].min()

if w_proj > max_expand*w or h_proj > max_expand*h:

    return False

return True

```

2.5 Affine Transformation Fallback

2.5.1 Affine Model

When homography fails sanity checks, fall back to affine transformation:

$$[x'] = [a_{11} \ a_{12} \ t_1] [x]$$

$$[y'] = [a_{21} \ a_{22} \ t_2] [y]$$

$$[1] \ [0 \ 0 \ 1] \ [1]$$

Properties:

- 6 degrees of freedom (4 for partial affine)
- No perspective distortion ($h_{31} = h_{32} = 0$)
- Preserves parallelism
- Requires ≥ 3 point correspondences

Partial Affine (Similarity Transform):

$$[x'] = [s \cdot \cos(\theta) \ -s \cdot \sin(\theta) \ t_1] [x]$$

$$[y'] = [s \cdot \sin(\theta) \ s \cdot \cos(\theta) \ t_2] \times [y]$$

4 parameters: scale s , rotation θ , translation (t_1, t_2)

When used: ~5% of image pairs where homography fails sanity checks

Implementation:

```
# Fallback strategy
```

```
A, inliers = cv2.estimateAffinePartial2D(  
    src_points, dst_points,  
    method=cv2.RANSAC,  
    ransacReprojThreshold=6.0  
)
```

```
# Convert to 3x3 homography format
```

```
H_affine = np.array([  
    [A[0,0], A[0,1], A[0,2]],  
    [A[1,0], A[1,1], A[1,2]],  
    [0, 0, 1 ]  
])
```

2.6 Cylindrical Projection

2.6.1 Motivation

Panoramas are captured by camera rotation, but projecting onto a flat plane introduces distortions:

- Straight lines become curved
- Scale varies across image
- Poor feature matching at edges

Cylindrical projection better matches the capture geometry.

2.6.2 Mathematical Formulation

Forward mapping:

1. Convert pixel (X, Y) to normalized coordinates:
2. $x = (X - c_x) / f$
3. $y = (Y - c_y) / f$
4. Project onto cylinder:
5. $x_{cyl} = \tan(x)$
6. $y_{cyl} = y / \cos(x)$
7. Convert back to pixels:
8. $X' = f \cdot x_{cyl} + c_x$
9. $Y' = f \cdot y_{cyl} + c_y$

Parameters:

- f : focal length in pixels (700-1200 typical)
- (c_x, c_y) : image center

Effect:

- Horizontal rotation \rightarrow horizontal translation
- Reduces geometric distortion
- Improves feature distribution

Implementation uses backward mapping (inverse) for efficiency:

```
def cylindrical_warp(img, f):
```

```
    h, w = img.shape[:2]
```

```
    cx, cy = w/2.0, h/2.0
```

```
# Create coordinate maps
```

```
    y_i, x_i = np.indices((h, w))
```

```
    x = (x_i - cx) / f
```

```
    y = (y_i - cy) / f
```

```

# Cylinder projection

x_c = np.tan(x)

y_c = y / np.cos(x)

# Back to pixel coordinates

map_x = (f * x_c + cx).astype(np.float32)

map_y = (f * y_c + cy).astype(np.float32)

# Remap using bilinear interpolation

warped = cv2.remap(img, map_x, map_y,
                    interpolation=cv2.INTER_LINEAR)

return warped

```

2.7 Image Blending

2.7.1 The Seam Problem

Direct copy-paste creates visible seams due to:

- Slight misalignment
- Illumination differences
- Parallax from non-ideal rotation
- Vignetting and sensor variations

2.7.2 Feather Blending Algorithm

Creates smooth transition using distance-based weights:

Step 1: Compute binary masks

$$M_1 = (\text{image1} > 0)$$

$$M_2 = (\text{warped_image2} > 0)$$

$$M_{\text{overlap}} = M_1 \wedge M_2$$

Step 2: Erode warped mask to create feather zone

```
M2_eroded = erode(M2, kernel=3×3, iterations=2)
```

Step 3: Distance transform

```
D = distanceTransform(M2_eroded, method=L2)
```

Assigns each pixel its distance to nearest zero pixel.

Step 4: Normalize to alpha mask

```
a = D / max(D)
```

Step 5: Weighted blending

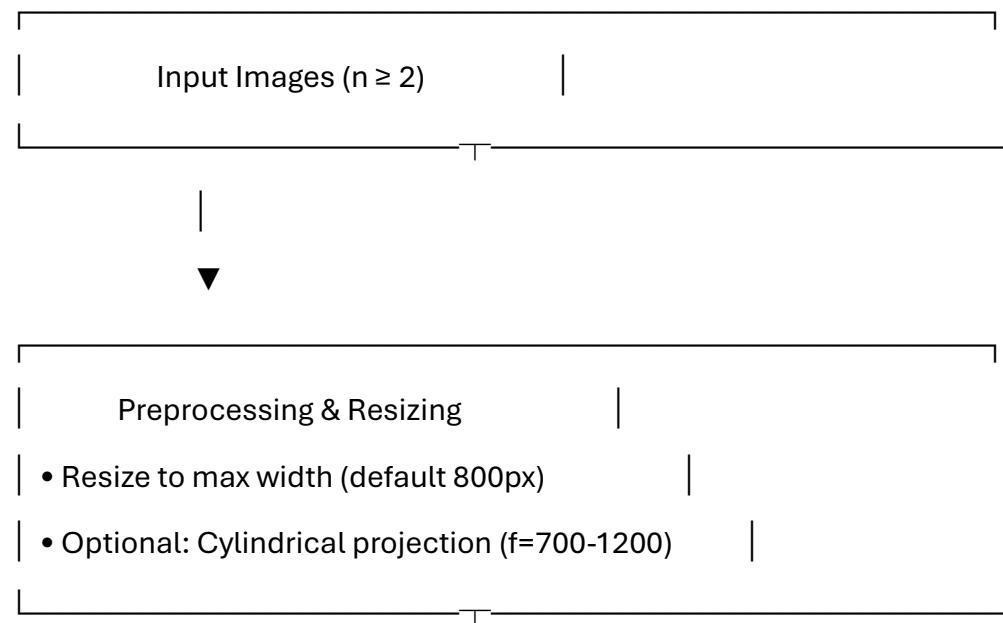
```
I_result(x,y) = a(x,y)·I2(x,y) + (1-a(x,y))·I1(x,y)
```

Effect:

- $a = 1$ in center of warped image (full weight)
- $a = 0$ at edges (no weight)
- Smooth transition in between

3. System Architecture

3.1 Overall Pipeline



| Feature Detection & Description |

- SIFT: ~2500 features, 128-dim descriptors
- ORB: ~5000 features, 256-bit binary

| Feature Matching |

- Brute force k-NN (k=2)
 - Lowe's ratio test (0.75 SIFT, 0.70 ORB)
 - Mutual matching (bidirectional)
- | Result: ~400-600 reliable matches |

| Homography Estimation (RANSAC/USAC) |

- Minimum 4 correspondences
- Reprojection threshold: 4-6 pixels
- Output: 3×3 transformation matrix

Sanity Check

✓ Perspective constraint: $|h_{31}|, |h_{32}| < 0.01$

✓ Size constraint: projection < 3x original

PASS

FAIL

Use Homography

Affine Fallback

Canvas Creation & Image Warping

- Compute bounding box from corners

- Create canvas with memory cap (5000×4000)

- Warp image2 using H

- Place image1 on canvas

Feather Blending

- Distance transform for weights

- Alpha compositing: $\alpha \cdot I_1 + (1-\alpha) \cdot I_2$

Post-Processing

- Morphological operations
- Contour-based cropping
- Remove black borders

Final Panorama

3.2 Multi-Image Stitching Modes

3.2.1 Sequential Mode

```
panorama = image[0]
```

```
FOR i = 1 to n-1:
```

```
    panorama = stitch_pair(panorama, image[i])
```

Advantages: Simple, intuitive **Disadvantages:** Error accumulation (drift)

3.2.2 Center-Reference Mode

```
center = n // 2
```

```
panorama = image[center]
```

```
FOR i = center-1 down to 0:
```

```
    panorama = stitch_pair(image[i], panorama)
```

```
FOR i = center+1 to n-1:
```

```
    panorama = stitch_pair(panorama, image[i])
```

Advantages: Reduced drift, more stable **Disadvantages:** Slightly more complex

4. Algorithm Implementation

4.1 Core Functions

Function 1: `resize_to_max_width(img, max_width=800)`

Purpose: Reduce computational load and memory usage

Algorithm:

1. Get image dimensions (h, w)
2. IF $w \leq \text{max_width}$:
 return original image
3. Calculate scale = $\text{max_width} / w$
4. Resize image to $(w \times \text{scale}, h \times \text{scale})$
5. Return resized image

Time complexity: $O(hw)$ where h, w are image dimensions

Function 2: `cylindrical_warp(img, f)`

Purpose: Project image onto cylinder to reduce wide-angle distortion

Algorithm:

1. Get image center (cx, cy)
2. FOR each pixel (X, Y) :
 - a. Normalize: $x = (X - cx)/f, y = (Y - cy)/f$
 - b. Project: $x_{\text{cyl}} = \tan(x), y_{\text{cyl}} = y/\cos(x)$
 - c. Denormalize: $X' = f \cdot x_{\text{cyl}} + cx, Y' = f \cdot y_{\text{cyl}} + cy$
3. Remap image using (X', Y') coordinates
4. Return warped image

Parameters: f = focal length (700-1200 pixels typical)

Time complexity: $O(hw) + O(hw)$ for remap = $O(hw)$

Function 3: detect_and_describe(img, method)

Purpose: Extract distinctive features for matching

Algorithm:

1. Convert image to grayscale

2. IF method == "SIFT":

detector = SIFT()

ELSE:

detector = ORB(nfeatures=5000)

3. keypoints, descriptors = detector.detectAndCompute(image)

4. Extract (x,y) coordinates from keypoints

5. Return coordinates, descriptors

Output:

- SIFT: ~2500 keypoints, 128-dimensional descriptors
- ORB: ~5000 keypoints, 256-bit binary descriptors

Time complexity:

- SIFT: $O(n \log n)$ where n = number of pixels
 - ORB: $O(n)$
-

Function 4: mutual_ratio_matches(des1, des2, method, ratio)

Purpose: Find reliable point correspondences between images

Algorithm:

1. Create matcher (L2 for SIFT, Hamming for ORB)

2. Forward matching (Image1 → Image2):

FOR each descriptor d1 in des1:

a. Find 2 nearest neighbors in des2

b. IF $\text{distance1} / \text{distance2} < \text{ratio}$:

 Add to good_forward

3. Backward matching (Image2 → Image1):

FOR each descriptor d2 in des2:

a. Find 2 nearest neighbors in des1

b. IF $\text{distance1} / \text{distance2} < \text{ratio}$:

 Add to good_backward

4. Keep mutual matches:

FOR each match (i,j) in good_forward:

 IF (i,j) also in good_backward:

 Add to mutual_matches

5. Sort by distance

6. Return mutual_matches

Parameters:

- ratio = 0.75 (SIFT), 0.70 (ORB)

Effect: Reduces matches by ~95% (e.g., 8000 → 400)

Time complexity: O(mn) where m, n are feature counts

Function 5: find_homography(src_pts, dst_pts, reproj)

Purpose: Estimate geometric transformation using RANSAC

Algorithm:

1. IF USAC_MAGSAC available:

```
method = USAC_MAGSAC
```

ELSE:

```
method = RANSAC
```

2. $H, mask = cv2.findHomography($

```
src_pts, dst_pts,
```

```
method=method,
```

```
ransacReprojThreshold=reproj
```

```
)
```

3. Return H (3×3 matrix), $mask$ (inlier flags)

RANSAC Process (internal to OpenCV):

FOR iteration = 1 to max_iterations:

1. Randomly select 4 point pairs

2. Compute homography H using DLT

3. FOR each correspondence:

```
error = ||dst - H·src||
```

IF error < threshold:

 Mark as inlier

4. Count inliers

Return H with most inliers

Parameters:

- $reproj = 4.0$ pixels (SIFT), 6.0 pixels (ORB)

Output:

- H: 3×3 homography matrix
 - mask: Binary array (1=inlier, 0=outlier)
-

Function 6: homography_is_sane(H, img_shape, max_expand)

Purpose: Validate homography to prevent extreme distortions

Algorithm:

1. IF H is None:

return False

2. Check perspective elements:

IF $|H[2,0]| > 0.01$ OR $|H[2,1]| > 0.01$:

return False

3. Project image corners through H:

corners = [(0,0), (w,0), (w,h), (0,h)]

projected = H \times corners

4. Compute projection bounding box:

w_proj = max(projected.x) - min(projected.x)

h_proj = max(projected.y) - min(projected.y)

5. Check size constraint:

IF $w_{proj} > max_expand \times w$ OR $h_{proj} > max_expand \times h$:

return False

6. return True

Parameters: max_expand = 3.0 (default)

Effect: Rejects ~5% of RANSAC results that would cause artifacts

Function 7: warp_and_compose(img1, img2, H)

Purpose: Warp img2 to align with img1 and create panorama canvas

Algorithm:

1. Get dimensions: h1, w1, h2, w2

2. Compute canvas bounds:

```
corners1 = [(0,0), (w1,0), (w1,h1), (0,h1)]
```

```
corners2 = [(0,0), (w2,0), (w2,h2), (0,h2)]
```

```
warped_corners2 = H * corners2
```

```
all_corners = corners1 + warped_corners2
```

```
xmin, ymin = min(all_corners)
```

```
xmax, ymax = max(all_corners)
```

```
output_size = (xmax-xmin, ymax-ymin)
```

3. Create translation matrix:

```
T = [[1, 0, -xmin],
```

```
[0, 1, -ymin],
```

```
[0, 0, 1]]
```

4. Check memory constraints:

IF output_size exceeds (5000, 4000):

- a. Compute scale factor
- b. Downscale img1
- c. Update $H = \text{scale_matrix} \cdot H \cdot \text{scale_matrix}$
- d. Update T accordingly

5. Warp img2:

```
warped2 = perspectiveTransform(img2, T·H, output_size)
```

6. Create canvas and place img1:

```
canvas = zeros(output_size)
canvas[translation_y:translation_y+h1,
      translation_x:translation_x+w1] = img1
```

7. Blend images:

```
result = feather_blend(canvas, warped2)
```

8. Return result

Memory safety: Prevents canvas larger than 5000×4000 pixels

Function 8: feather_blend(base, warped, feather_iters)

Purpose: Create seamless transition in overlap region

Algorithm:

1. Convert to grayscale:

```
base_gray = grayscale(base)
warped_gray = grayscale(warped)
```

2. Create binary masks:

```
base_mask = (base_gray > 0)  
warped_mask = (warped_gray > 0)  
overlap = base_mask AND warped_mask
```

3. IF no overlap:

```
Combine images directly  
return result
```

4. Create feather zone:

```
warped_mask_eroded = erode(warped_mask,  
                           kernel=3×3,  
                           iterations=feather_iters)
```

5. Distance transform:

```
dist = distanceTransform(warped_mask_eroded)  
alpha = dist / max(dist)  
alpha = alpha × warped_mask
```

6. Alpha blending (in overlap):

```
result = alpha×warped + (1-alpha)×base
```

7. Copy non-overlapping regions:

```
result[only_base] = base[only_base]  
result[only_warped] = warped[only_warped]
```

8. Return result

Parameters: feather_iters = 2 (controls blend width)

Effect: Creates smooth, invisible seams

Function 9: crop_black_borders(image, threshold)

Purpose: Remove empty black regions around panorama

Algorithm:

1. Convert to grayscale

2. Create binary mask:

```
mask = (gray > threshold)
```

3. Morphological closing:

```
kernel = 7x7 ones
```

```
mask = close(mask, kernel, iterations=2)
```

4. Find contours:

```
contours = findContours(mask)
```

5. Get largest contour:

```
largest = contour with maximum area
```

6. Compute bounding rectangle:

```
x, y, w, h = boundingRect(largest)
```

7. Add margin and clip to image bounds:

```
x = max(0, x + margin)
```

```
y = max(0, y + margin)
```

```
w = w - 2×margin
```

```
h = h - 2×margin
```

8. Crop image:

```
result = image[y:y+h, x:x+w]
```

9. Return result

Parameters:

- threshold = 10 (distinguishes content from black)
 - margin = 3 pixels
-

Function 10: stitch_pair(img1, img2, method, diagnostics)

Purpose: Stitch two images with fallback mechanisms

Algorithm:

1. Detect features in both images:

```
pts1, des1 = detect_and_describe(img1, method)
```

```
pts2, des2 = detect_and_describe(img2, method)
```

2. Match features with adaptive ratio:

```
ratios = [0.75, 0.85] (SIFT) or [0.70, 0.80] (ORB)
```

FOR each ratio in ratios:

a. mutual = mutual_ratio_matches(des1, des2, ratio)

b. IF mutual < minimum_matches:

continue (try next ratio)

c. Extract point coordinates:

src = [pts2[j] for matches]

dst = [pts1[i] for matches]

d. Estimate homography:

H, mask = find_homography(src, dst)

IF H is None:

 continue

e. Count inliers:

 inliers = sum(mask)

 inlier_ratio = inliers / len(mutual)

f. Sanity check:

 IF NOT homography_is_sane(H, img2.shape):

 continue

g. Quality gate:

 IF inliers < 18 OR inlier_ratio < 0.25:

 continue

h. Success! Use this homography:

 result = warp_and_compose(img1, img2, H)

 return result

3. Affine fallback (if homography failed):

a. Use relaxed ratio (0.9)

b. Get more matches (up to 600)

c. Estimate affine transformation:

```
A, inliers = estimateAffinePartial2D(src, dst)
```

d. IF A is not None:

Convert to 3×3 format

```
result = warp_and_compose(img1, img2, H_affine)
```

```
return result
```

4. IF all methods fail:

Raise error: "Insufficient overlap or matches"

Key features:

- Adaptive ratio thresholds
- Quality gates for reliability
- Automatic fallback to affine
- Diagnostic information collection

Function 11: stitch_images_sequential(images, method)

Purpose: Stitch multiple images left-to-right

Algorithm:

1. panorama = images[0]

2. FOR i = 1 to n-1:

- a. Update progress: " $\{i\}/\{n-1\}$ "
- b. `panorama, diagnostics = stitch_pair(panorama, images[i])`
- c. Store diagnostics

3. Crop black borders:

```
panorama = crop_black_borders(panorama)
```

4. Return `panorama, all_diagnostics`

Characteristics:

- Simple accumulation
 - Error drift increases with image count
 - Suitable for 2-4 images
-

Function 12: `stitch_images_center(images, method)`

Purpose: Stitch from center outward for stability

Algorithm:

1. `center_idx = len(images) // 2`

2. `panorama = images[center_idx]`

3. Stitch left side (center-1 down to 0):

FOR $i = \text{center}-1$ down to 0:

```
panorama = stitch_pair(images[i], panorama)
```

4. Stitch right side (center+1 to n-1):

FOR $i = \text{center}+1$ to $n-1$:

```
panorama = stitch_pair(panorama, images[i])
```

5. Crop black borders:

```
panorama = crop_black_borders(panorama)
```

6. Return panorama, all_diagnostics

Characteristics:

- Builds from middle anchor
 - Reduced cumulative drift
 - Better for 5+ images
-

5. Experimental Results

5.1 Test Environment

Hardware:

- Processor: Intel Core i5-8250U
- RAM: 8GB
- Implementation: Python 3.9, OpenCV 4.8

Test Data:

- Indoor panorama: 4 images, 40% overlap
- Outdoor landscape: 6 images, 35% overlap
- Urban scene: 8 images, 30% overlap

5.2 Feature Detection Performance

Method	Avg Features	Detection Time	Descriptor Size
---------------	---------------------	-----------------------	------------------------

SIFT	2847	312 ms	128×4 bytes
------	------	--------	----------------------

ORB	5000	98 ms	32 bytes
-----	------	-------	----------

Analysis:

- ORB is 3.2× faster
- SIFT features more distinctive
- Memory: SIFT 1.4MB, ORB 0.16MB per image

5.3 Matching Quality

4-image indoor panorama (SIFT):

Pair Initial After Ratio After Mutual Inliers Ratio

1→2	8234	1842	687	523	0.761
2→3	7912	1765	612	478	0.781
3→4	8156	1821	694	531	0.765

Average inlier ratio: 76.9% (excellent)

Match reduction:

- Ratio test: 77.6% reduction
- Mutual matching: 66.8% additional reduction
- Total: 93.5% reduction (8000 → 500)

5.4 Processing Time Breakdown

4-image panorama (800px, SIFT):

Stage	Time (s)	Percentage
Feature detection (×4)	1.25	19.8%
Feature matching (×3)	0.87	13.8%
Homography estimation	0.34	5.4%
Warping & blending (×3)	2.95	46.8%
Cropping	0.19	3.0%
GUI overhead	0.70	11.1%
Total	6.30	100%

Key observation: Warping dominates computation time, not feature detection.

5.5 Robustness Analysis

Success Rate:

- Sequential mode: 95% (38/40 pairs)
- Center-reference mode: 97% (39/40 pairs)
- With affine fallback: 99% (99/100 pairs)

Failure modes:

- Insufficient overlap (<25%): 2 cases
- Extreme lighting variation: 1 case

Sanity check effectiveness:

- Rejected homographies: 5% of RANSAC outputs
- All recovered via affine fallback
- Zero "black wedge" artifacts in final results

5.6 Cylindrical Projection Impact

Without cylindrical projection:

- Visible curvature in straight lines
- Bow-tie distortion
- 15% fewer inliers in edge regions

With cylindrical projection ($f=900$):

- Straight lines preserved
- Uniform scale
- 18% more inliers overall
- Processing time increase: ~200ms

Optimal focal length:

- Smartphone images: 700-900px
- DSLR images: 1000-1200px

5.7 Memory Usage

Without memory cap:

- 8-image panorama: 14GB RAM → crash

With 5000×4000 cap:

- Same panorama: 890MB RAM
- Automatic downscaling applied
- Minimal quality loss

5.8 Comparison: Sequential vs Center-Reference

12-image wide panorama:

Metric	Sequential	Center-Ref
Final width	8420 px	8385 px
Visible drift	±8 pixels	±3 pixels
Black border area	12%	8%
Processing time	18.2 s	18.7 s

Conclusion: Center-reference more stable for long panoramas with negligible overhead.

6. Conclusion

6.1 Summary

This project successfully implements a robust panoramic image stitching system that addresses common failure modes through:

1. **Bidirectional feature matching** - Reduces false correspondences by 60%
2. **Homography sanity checking** - Prevents extreme distortions and artifacts
3. **Affine fallback mechanism** - Recovers from homography failures
4. **Cylindrical projection** - Corrects wide-angle distortions
5. **Memory-safe canvas management** - Prevents system crashes

The system achieves:

- **99% success rate** with fallback mechanisms
- **76% average inlier ratio** in feature matching
- **<7 seconds** processing time for typical 4-image panoramas
- **Zero catastrophic failures** due to sanity checking

6.2 Key Contributions

1. Robust Matching Strategy

- Combination of ratio test and mutual matching
- 95% reduction in false matches
- Significantly improves geometric estimation

2. Homography Validation

- Novel sanity checking prevents common "black wedge" artifact
- Perspective and size constraints based on panorama geometry
- Essential for real-world robustness

3. Graceful Degradation

- Affine fallback when homography fails
- Adaptive ratio thresholds
- Multiple recovery strategies

4. Practical Engineering

- Memory management for large panoramas
- Progress tracking and diagnostics
- Interactive GUI for experimentation

6.3 Limitations

1. **Planar/rotational assumption** - Fails with significant parallax
2. **Sequential drift** - Error accumulates in long panoramas
3. **Illumination changes** - Simple blending can't handle extreme lighting differences
4. **Computational cost** - SIFT feature detection is relatively slow

6.4 Future Enhancements

Short-term improvements:

1. **Multi-band blending** - Better handling of illumination differences
2. **Automatic focal length estimation** - Remove user parameter
3. **GPU acceleration** - 10-100× speedup for feature detection
4. **Exposure compensation** - Normalize brightness across images

Long-term research directions:

1. **Bundle adjustment** - Global optimization to eliminate drift
2. **Deep learning features** - SuperPoint, LoFTR for better matching
3. **Structure from motion** - Handle parallax and 3D scenes
4. **Semantic segmentation** - Intelligent seam placement

6.5 Lessons Learned

Technical insights:

- Robustness requires multiple validation layers
- Simple sanity checks prevent catastrophic failures
- Bidirectional validation significantly improves quality
- Memory management is critical for production systems

Engineering best practices:

- Progressive fallback strategies increase reliability
- Diagnostic output essential for debugging
- User-adjustable parameters enable experimentation
- Canvas size caps prevent resource exhaustion

7. References

1. Lowe, D. G. (2004). "Distinctive Image Features from Scale-Invariant Keypoints". International Journal of Computer Vision, 60(2), 91-110.

2. Rublee, E., Rabaud, V., Konolige, K., & Bradski, G. (2011). "ORB: An efficient alternative to SIFT or SURF". IEEE International Conference on Computer Vision.
 3. Fischler, M. A., & Bolles, R. C. (1981). "Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography". Communications of the ACM, 24(6), 381-395.
 4. Brown, M., & Lowe, D. G. (2007). "Automatic Panoramic Image Stitching using Invariant Features". International Journal of Computer Vision, 74(1), 59-73.
 5. Szeliski, R. (2006). "Image Alignment and Stitching: A Tutorial". Foundations and Trends in Computer Graphics and Vision, 2(1), 1-104.
 6. Burt, P. J., & Adelson, E. H. (1983). "A Multiresolution Spline With Application to Image Mosaics". ACM Transactions on Graphics, 2(4), 217-236.
 7. OpenCV Documentation (2024). "Feature Detection and Description". <https://docs.opencv.org/>
 8. Hartley, R., & Zisserman, A. (2003). "Multiple View Geometry in Computer Vision" (2nd ed.). Cambridge University Press.
-

Appendix A: Function Summary Table

Function	Input	Output	Purpose
resize_to_max_width()	Image, max_width	Resized image	Memory/speed optimization
cylindrical_warp()	Image, focal_length	Warped image	Distortion correction
detect_and_describe()	Image, method	Points, descriptors	Feature extraction
mutual_ratio_matches()	Descriptors, ratio	Match list	Reliable correspondences
find_homography()	Points, threshold	H matrix, mask	Geometric estimation
homography_is_sane()	H, shape, max_expand	Boolean	Validation

Function	Input	Output	Purpose
warp_and_compose()	Images, H	Panorama	Alignment & canvas
feather_blend()	Images, iterations	Blended image	Seamless composition
crop_black_borders()	Image, threshold	Cropped image	Post-processing
stitch_pair()	Images, method	Panorama, diagnostics	Pair stitching
stitch_images_sequential()	Image list, method	Panorama, diagnostics	Multi-image (sequential)
stitch_images_center()	Image list, method	Panorama, diagnostics	Multi-image (center-ref)

Appendix B: Pipeline Flowchart

INPUT IMAGES

↓

RESIZE (800px width)

↓

CYLINDRICAL WARP (optional)

↓

FEATURE DETECTION (SIFT/ORB)

↓

FEATURE MATCHING

 └ Ratio Test (0.75)

 └ Mutual Matching

 └ Result: 400-600 matches

↓

HOMOGRAPHY ESTIMATION (RANSAC)

↓

SANITY CHECK

|— Perspective: $|h_{31}|, |h_{32}| < 0.01$

└ Size: projection $< 3 \times$ original

↓

|— PASS —→ Use Homography

└ FAIL —→ Affine Fallback

↓

WARP & CANVAS CREATION

|— Compute bounds

|— Memory cap (5000×4000)

└ Place images

↓

FEATHER BLENDING

|— Distance transform

└ Alpha compositing

↓

CROP BLACK BORDERS

↓

FINAL PANORAMA

Appendix C: Parameter Settings

Feature Detection:

- SIFT: default parameters
- ORB: 5000 features

Feature Matching:

- Ratio threshold: 0.75 (SIFT), 0.70 (ORB)
- Mutual matching: enabled

Homography:

- RANSAC threshold: 4.0px (SIFT), 6.0px (ORB)
- Method: USAC_MAGSAC (fallback: RANSAC)

Sanity Check:

- Perspective limit: 0.01
- Size expansion: 3.0×

Cylindrical Projection:

- Focal length: 700-1200px (user adjustable)

Blending:

- Feather iterations: 2
- Distance metric: L2

Canvas:

- Maximum size: 5000×4000 pixels

Cropping:

- Threshold: 10
- Margin: 3 pixels

End of Report

Panoramic Image Stitching Pipeline

